

Agentic AI Final Project Milestone II: Agentic System Risks

RISKS	Mitigation Plan	
	Design & Development of agent	Agent Integration into process
Adversarial attacks	Validate user inputs (grades, attendance, etc.) and filter prompts to prevent manipulation of the agent's reasoning (e.g., fake academic data or prompt injection).	Monitor unusual input patterns and restrict excessive or repeated submissions; implement user authentication for access.
Bias and Discrimination	Ensure the model does not make assumptions based on non-academic factors (e.g., background, demographics) and test outputs across diverse student scenarios.	Allow users to review and challenge recommendations; include transparency in how risk levels are determined.
Misinformation & Hallucinations	Constrain outputs to academic performance data only (grades, attendance, assignments) and use structured prompts to reduce hallucinations.	Include disclaimers that recommendations are advisory and encourage consultation with academic advisors for final decisions.
Data Privacy Violations	Avoid storing personally identifiable student data; anonymize inputs and limit data collection to necessary academic metrics.	Secure user inputs and responses, and ensure compliance with student data privacy standards (e.g., FERPA-like principles).
Intellectual Property Risks	Ensure training and prompts do not use copyrighted academic materials without permission.	Prevent the system from generating or reproducing proprietary course content or assignments.