# Project Proposal: AI Terms & Policy Translator

1. **Problem Statement & Motivation**

In today's digital world, users are constantly asked to agree to terms and conditions, privacy policies, and usage agreements when applying for internships, using software, or signing up for online platforms. These documents are usually very long and written in complex legal language that most people do not fully understand. Because of this, many users simply click "I agree" without reading the terms so they can move forward quickly. I personally do this often, especially when applying for internships or downloading software, because the policies feel overwhelming and time-consuming to read.

Later on, this behavior can lead to unexpected problems such as personal data being shared, automatic subscription renewals, or limits on user rights that the user was not aware of. This is a real-world business and consumer problem. Users may feel frustrated or misled after discovering issues hidden in the terms, while companies may face trust and reputation concerns. The main issue is not that policies are unavailable, but that they are difficult for the average user to understand.

The goal of this project is to design an AI agentic system that translates policies and terms into plain English summaries and highlights potential risks so users, including people like me, can better understand what they are agreeing to before clicking "I agree.

The target users for this system include college students applying for internships or jobs, everyday consumers signing up for apps or software tools, and non-technical users who do not have legal or policy knowledge. Many of these users, myself included, want to move quickly through sign-up processes but still want to understand important details that could affect privacy, finances, or future access to services.

An agentic AI solution is appropriate because understanding policies requires multiple steps that benefit from specialization. A human could read and interpret policies, but this takes time and effort, and most users are not willing or able to do this every time they encounter a new agreement. A single AI response is also limited because policy understanding involves parsing document structure, simplifying language, and identifying risks. By using multiple AI agents with specific roles, the system can break the task into clear steps. One agent focuses on parsing the document, another simplifies the language, and another highlights potential risks. This approach better reflects how complex documents are analyzed and provides clearer, more reliable results than a single automated summary.

2. **Use Cases & Scenarios**

**Scenario 1: Internship Application Platform Terms**

As a college student applying for internships, I often encounter online platforms that require me to agree to long terms and conditions and privacy policies before submitting my application. I usually click “I agree” without fully reading them because I do not understand most of the language and want to continue with the application. In this scenario, the student uploads the policy into the AI Policy & Terms Translator. First, the Legal / Policy Parser Agent scans the document and identifies sections related to data collection, data sharing, and user consent. It breaks the document into organized sections for easier analysis. Next, the Simplification Agent rewrites these sections in plain English, explaining what the policy actually means for the student, such as how their resume and personal information may be used. Finally, the Risk Highlight Agent flags any concerning clauses, such as sharing data with third parties or storing data after the application process ends. This helps the student decide whether to continue while being more informed.

**Scenario 2: Software Subscription and Auto-Renewal Agreement**

When signing up for software tools for school or work, I often see subscription agreements that mention billing and renewal terms that I usually do not read carefully. This can be a problem if the subscription renews automatically or has strict cancellation rules. In this scenario, the user uploads the subscription agreement into the system. The Legal / Policy Parser Agent identifies sections related to pricing, billing cycles, auto renewal, and cancellation policies. The Simplification Agent then translates these sections into clear, simple language explaining how much the user will be charged, how often payments occur, and how cancellation works. The Risk Highlight Agent flags clauses such as automatic renewals or non-refundable payments so the user understands potential financial risks before agreeing.

## 3. Initial System Design

The system will include a few different agents that each focus on a specific task. One agent will act as a legal or policy parser and will break long policy documents into smaller sections while identifying important topics like privacy, billing, and termination. Another agent will focus on simplifying the language by rewriting complex legal terms into plain English that is easier for users like me to understand. The final agent will look for potentially risky clauses and highlight them with short explanations so users can quickly see what parts of the policy might be concerning before agreeing to it. These agents will operate in a multi-agent workflow, where the output of one agent is passed to the next. Right now, I am not completely sure which tools or APIs I will use for this project. The main idea is that users will be able to upload or paste policies, and the system will analyze the text and explain it in simpler language. I expect to use basic tools that allow the system to read text or PDF files and an AI model to help summarize and explain the policies. As I continue working on the project, I will explore different tools to see what works best and is easiest to use. The final tools will be chosen based on what is realistic to implement and fits the project requirements.

## 4. Feasibility & Risks

Potential challenges include misinterpreting legal language, handling very long documents, and oversimplifying important details. To reduce these risks, the system will clearly state that it provides informational summaries and not legal advice. Privacy concerns will be addressed by limiting document storage and processing documents only as needed. The project is feasible within the course timeline and computing constraints. The project will be completed within the duration of the course. Early weeks will focus on finalizing the system design and building the front-end interface. The middle of the course will focus on implementing the multi-agent workflow. The final weeks will be used for evaluation, testing, and preparing the final report and presentation.